

# Lab 1

## 3.1

- c. The ancestor process runs in an infinite loop for the rest of its existence.
- d. `Nulluser()` does not ever return from `start.S`.
- e. `Halt()` is in `intr.S` and it stops the CPU.
- f. When I removed the `halt()` call after `nulluser()`, nothing changed.
- g. When I replaced the while loop at the end of `nulluser()` with a call to `halt()`, nothing changed.

## 3.2

- a. After a new process is created in Linux using `fork()`, the parent and child processes both run.
- b. Since a new process has been made obviously with the intent to run, I believe the child process will run before the parent.
- c. As an app programmer my preference would be that the child process runs first. This is because the parent process may rely on a result from the child process, so that executing first would be optimal.
- e. The difference between `newProcess()` and `create()` is that `create()` makes a process from nothing, while `newProcess()` duplicates the parent process.
- f. The `create()` method is different than `clone()` in that the process made by `create()` has its own execution context while the process made by `clone()` can share it with the parent process.
- g. The `posix_spawn()` function differs from `newProcess()` in that the new process is created in an image passed in as an argument while `newProcess()` uses the parent process image to create a child process.
- h. There is no best way to create processes because every context is different when creating a process. Different situations may be optimized by a specific kind of process creation, but considering the amount of processes and how quick they all are in general there is no best way to create a process.